

# ECE 454/750: Distributed Computing

## Tutorial 3: Assignment 1

TA: Hua Fan

[h27fan@uwaterloo.ca](mailto:h27fan@uwaterloo.ca)

May 26, 2015 05:30-06:20, QNC 2502

# Outline

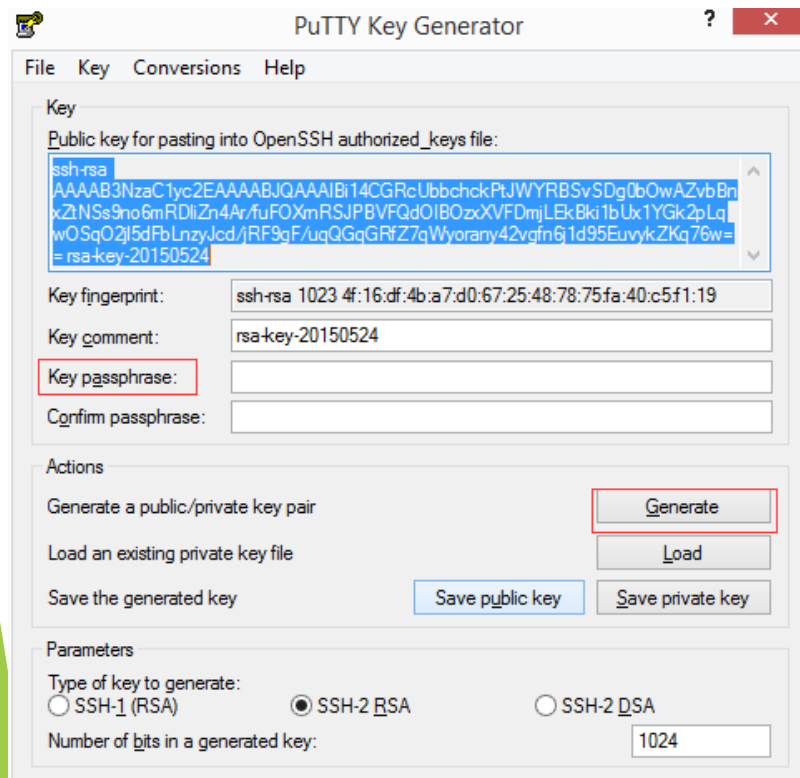
In this tutorial we will discuss:

- ▶ Public Key Authentication (windows, linux)
- ▶ Set CPU affinity
- ▶ Submission Instructions
- ▶ Design document
- ▶ Build Instructions
- ▶ Sample test case
- ▶ Q & A
- ▶ Slides will be uploaded in Learn later.
- ▶ Discussion forum:  
[piazza.com/uwaterloo.ca/summer2015/ece454750](https://piazza.com/uwaterloo.ca/summer2015/ece454750)

# Public Key Authentication With PuTTY\*

- ▶ No need to send passwd to the server you want to ssh.
- ▶ Send public key to remote server; the local machine stores both the public key and the private
- ▶ Step 1: generate public/private key pair by PuTTYgen
  - ▶ Click Generate button
  - ▶ Save private key by button
  - ▶ Save public key in blue block in the following picture(highlight with the mouse)

# Step 1: generate public/private key pair



**PuTTY Key Generator**

File Key Conversions Help

**Key**

Public key for pasting into OpenSSH authorized\_keys file:

```
ssh-rsa AAAAB3NzaC1yc2EAAAABJQAAAIBi14CGRcUbbchckPtJWYRBSySDg0bOwAZybBn  
xZiNSs9no6mRDliZn4Ar/fuFOXmRSJPBVFQdOIBOzcXVFDmjLEkBki1bUx1YGk2pLq  
wOSqO2j5dFbLnzyJcd/jRF9gF/uqQGqGrfZ7qWyorany42vgfn6j1d95EuvykZKq76w=  
=rsa-key-20150524
```

Key fingerprint: ssh-rsa 1023 4f:16:df:4b:a7:d0:67:25:48:78:75:fa:40:c5:f1:19

Key comment: rsa-key-20150524

Key passphrase:

Confirm passphrase:

**Actions**

Generate a public/private key pair **Generate**

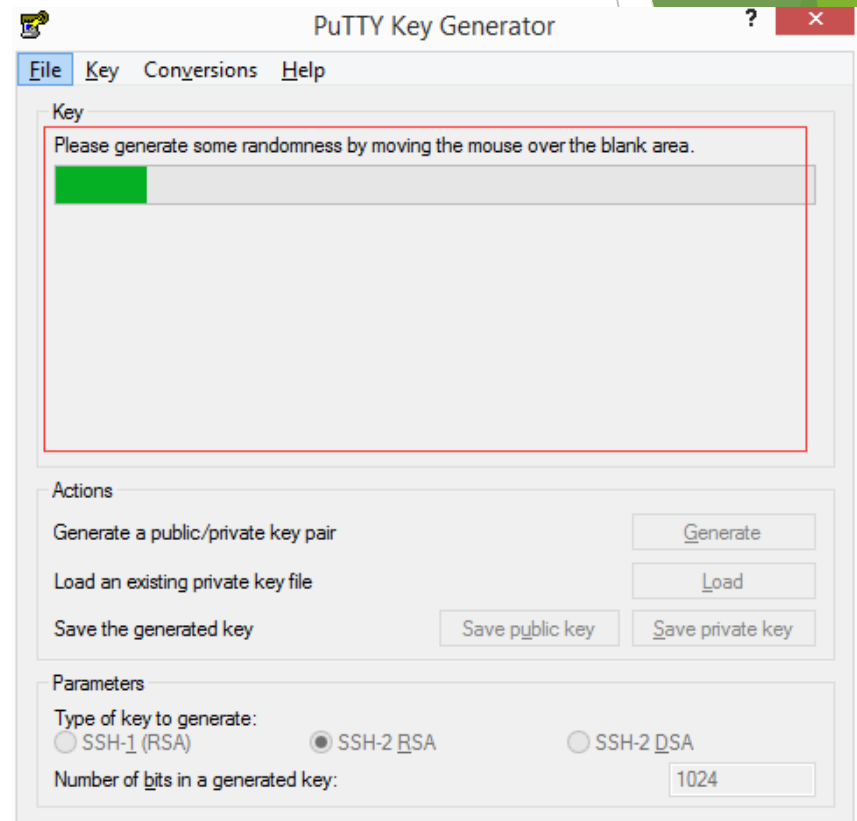
Load an existing private key file **Load**

Save the generated key **Save public key** **Save private key**

**Parameters**

Type of key to generate:  
☐ SSH-1 (RSA) ☒ SSH-2 RSA ☐ SSH-2 DSA

Number of bits in a generated key: 1024



**PuTTY Key Generator**

File Key Conversions Help

**Key**

Please generate some randomness by moving the mouse over the blank area.

**Actions**

Generate a public/private key pair **Generate**

Load an existing private key file **Load**

Save the generated key **Save public key** **Save private key**

**Parameters**

Type of key to generate:  
☐ SSH-1 (RSA) ☒ SSH-2 RSA ☐ SSH-2 DSA

Number of bits in a generated key: 1024

## Step 2: install the public key on the remote host

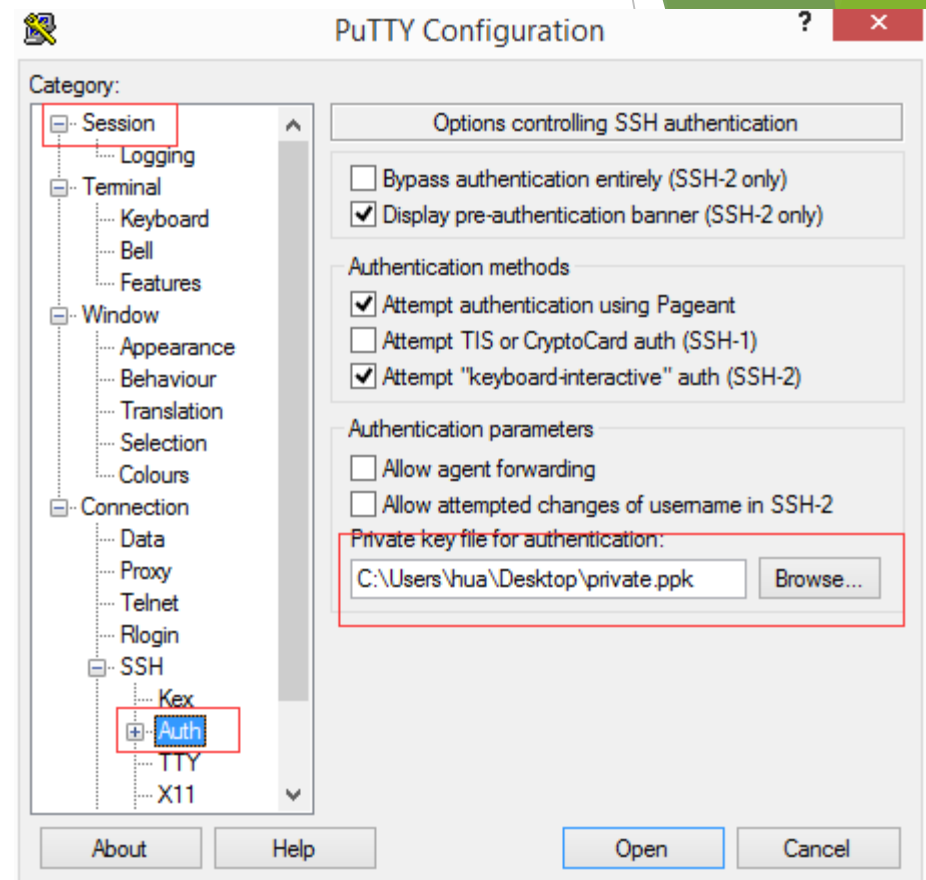
- ▶ Append the context of public key to file  
~/.ssh/authorized\_keys on the remote host(linux)
- ▶ One way: edit that file by text editor

```
3 ssh-rsa AAAAB3NzaC1yc2EAAAABJQAAIB7jdaNjrOSzIGrBhYLu8
lYIrP+r30PwGQjHM9/P5YhBohndqUlsQixjrOO7ZXAETZzNduDpmtR
3oy4Tryhd2PQpPhwGirw84I6zP2Scr8teu/DzUY3apZ1K8G7s1+Ydq
B6c1PVJJC/TswHLGCYx610Mokwwvhf8+GUSb+8tKFZ9Q= rsa-key
-20150524
```

- ▶ Another way: transfer the public key file to remote host
  - ▶ `cat public.key >> ~/.ssh/authorized_keys`

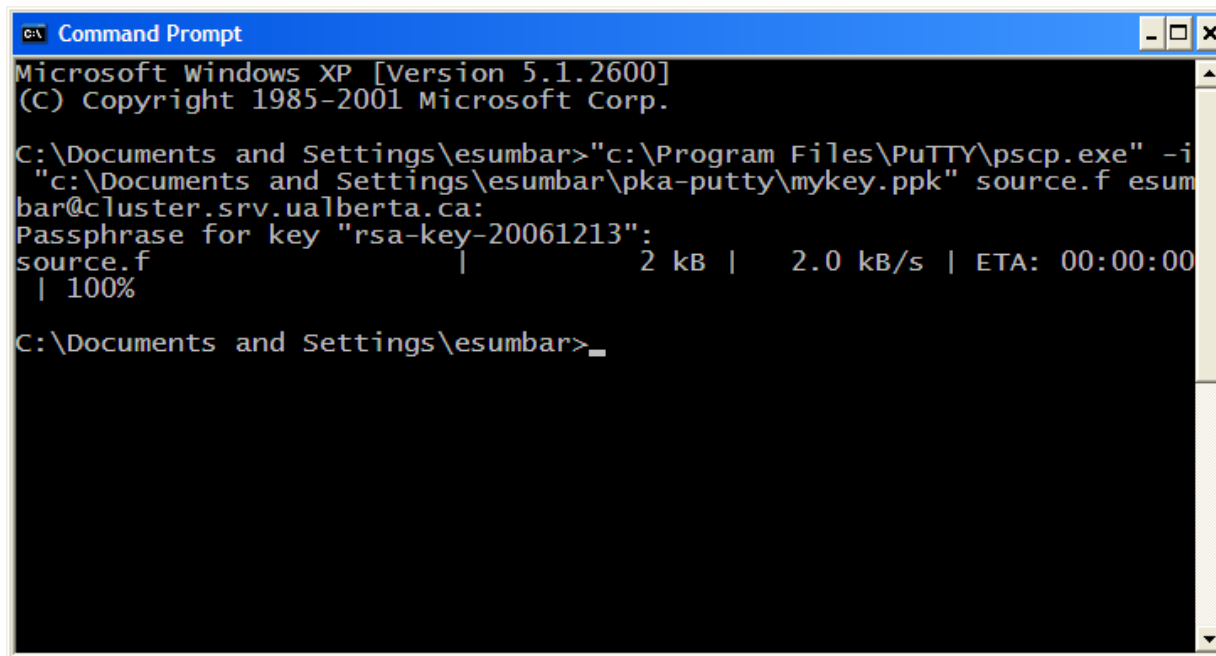
# Step 3: Setting private key for a session

- ▶ Connection->SSH->Auth:
  - ▶ Select the private key file.
- ▶ Save session: session->Save



# Additional usage

- ▶ File transfers with key authentication
  - ▶ `pscp.exe -i` flag to specify private key.
  - ▶ Usage(similar to scp on linux): `pscp.exe -i private.key file user@host:path`



```
C:\ Command Prompt
Microsoft Windows XP [Version 5.1.2600]
(C) Copyright 1985-2001 Microsoft Corp.

C:\Documents and Settings\esumbar>"c:\Program Files\PuTTY\pscp.exe" -i
"c:\Documents and Settings\esumbar\pka-putty\mykey.ppk" source.f esum
bar@cluster.srv.ualberta.ca:
Passphrase for key "rsa-key-20061213":
source.f                               2 kB |  2.0 kB/s | ETA: 00:00:00
| 100%

C:\Documents and Settings\esumbar>_
```

\* Figure from: <http://www.ualberta.ca/CNS/RESEARCH/LinuxClusters/pka-putty.html>

# Additional usage

- ▶ PuTTY authentication agent, pageant
  - ▶ If you don't want type passphrase every time.
  - ▶ Right-click on the icon and choose "Add Key" as illustrated below.





# Public Key Authentication With ssh-keygen (Linux)

- ▶ Step 1: generate key pair, execute: `ssh-keygen`

```
[h27fan@eceubuntu ~]$ ssh-keygen
Generating public/private rsa key pair.
Enter file in which to save the key (/home/h27fan/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Passphrases do not match. Try again.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in ./test-keygen.
Your public key has been saved in ./test-keygen.pub.
The key fingerprint is:
36:91:1f:b8:49:29:d0:59:4c:51:87:a0:35:f4:70:b4 h27fan@eceubuntu.
The key's randomart image is:
+---[ RSA 2048]-----+
  .. *X=+..
  .00.O.O
  O * E
  O = .
  S .
  . .
+-----+
```

# Linux ssh authen. (cont.)

- ▶ Step 2: install public key to remote host  
~/.ssh/authorized\_keys (the same as before)
- ▶ Step 3: Place private key in local machine ~/.ssh/
  - ▶ If use the default path to save keys, there on need for this step
- ▶ Step 4: chmod 600, so no other users can read them
  - ▶ `chmod 600 ~/.ssh/id_rsa*`
  - ▶ `chmod 600 ~/.ssh/authorized_keys`

# Set CPU affinity by taskset

- ▶ Why use multi-core machines?
  - ▶ Parallel execution on multi-cores by multi-threading programming
  - ▶ Physic speed limit of CPU cores.
  - ▶ Hardware support: eceubuntu 6 cores; ecelinux1 8 cores
- ▶ Why need CPU affinity?
  - ▶ Reduce cache miss
  - ▶ Establish hardware resource bound
- ▶ Taskset
  - ▶ *taskset -c cpu-list <application>*
  - ▶ e.g. *taskset -c 0,1,2,3 <application>*

# Submission Instructions

- ▶ You should place all your files in one folder, called **a1ece454750**. The folder should contain:
  - ▶ build.xml
  - ▶ design document in pdf file.
  - ▶ All source code(java, thrift files) organized at your own choice, but files built by an **ant** command.
- ▶ You should tar and gzip the folder by:
  - ▶ **tar czf a1ece454750.tar.gz a1ece454750**
  - ▶ we will untar and gunzip by: **tar xzf a1ece454750.tar.gz**
- ▶ Don't submit any .class files or gen-java/\*.java files, which can be generated by your source code.

# Design document

- ▶ Short : 1-page
- ▶ The reasoning of your choice of sync/async RPCs, multi-threading servers, and protocols.
- ▶ Please be consistent with you choice on all function implementations(sync/async, protocols), otherwise it will be difficult for us to test.
- ▶ How do you achieve load balancing?
- ▶ How do you detect and deal with crash failures?
- ▶ How do you ensure that a BE node that joins the cluster can receive requests from the FE layer within 1 second of startup?

# Build Instructions

- ▶ cd a1ece454750, where build.xml placed in, and build your solution by executing: **ant**
- ▶ Build must jar the class files and auxiliary resources and output as **ece454750s15a1.jar**. The jar file must place in the folder a1ece454750. Refer how tutorial.jar is generated in tutorial 1&2.
- ▶ You can safely have these assumptions:
  - ▶ Build on **eceubuntu**
  - ▶ Thrift compiler 0.9.1 located in: **/usr/bin/thrift**
  - ▶ There is a folder, **lib**, sharing the same parent folder with a1ece454750. The lib folder (referred as ../lib/ in build.xml) contains:
    - ▶ libthrift-0.9.1.jar
    - ▶ jbcrypt.jar
    - ▶ commons-codec-1.6.jar commons-logging-1.1.1.jar httpcore-4.2.4.jar log4j-1.2.14.jar slf4j-api-1.5.8.jar commons-lang3-3.1.jar httpclient-4.2.5.jar junit-4.4.jar servlet-api-2.5.jar slf4j-log4j12-1.5.8.jar (from thrift-0.9.1/lib/java/build/lib)
    - ▶ Other libs/source code are not allowed unless approved by course instructor and TA.

# Test runnable

- ▶ Execute under folder a1ece454750 : `cd a1ece454750`
- ▶ Test on both ecelinux and eceubuntu. Instructions for cross-compiling to java 1.6 were given in tutorial 1.
- ▶ `java -cp "ece454750s15a1.jar:../lib/*" ece454750s15a1.FEServer -host ecelinux1 -pport 8123 -mport 9123 -ncores 2 -seeds ecelinux1:10123,ecelinux2:10123,ecelinux3:10123`
- ▶ The port number for the seed corresponds to the management interface.

# Some notes on testing

- ▶ We will use our own client for testing, and this client will know the host names and port numbers of the FE nodes. In other words, students do **not** have to implement a mechanism by which clients would discover the FE nodes.
- ▶ We will implement different kinds of clients and detect your RPC mode(sync or async) and protocol choice by them.
- ▶ Some FE nodes will be started as seeds, meaning that **as seed**: the host and management port of the FE node appear in the list of seed nodes.
- ▶ Every FE and BE node is given **the same** list of seed nodes.



# Sample test case

## ► Basic:

- 1. Start a FE as seed on ecelinuxX taking one core
- 2. Start a BE on ecelinuxX taking another core
- 3. Test each interface, **expecting** results as specification.
- 4. Start another non-seed FE “on the fly” on ecelinuxY
- 5. sleep 1 second, new FE should be able to forward requests to the BE node.
- 6. Test step 3 again against FE on ecelinuxY

## ► Note

- X,Y stands for 1, 2, 3, 5
- Test on servers as early as possible. **Don't leave it to the last day.** There are many students in the class and the infrastructure will become very busy shortly before the deadline.
- Avoid testing your code using eceubuntu to prevent overload, use ecelinux machines instead.

# Sample: balanced load

- ▶ Start a FE as seed on ecelinux1 taking 4 cores
- ▶ Start a BE on ecelinux2 taking 2 cores
- ▶ Start another BE on ecelinux2 taking 2 cores
- ▶ Client send requests
- ▶ Request performance counters on the servers
- ▶ **Expecting:** load is roughly split in half between the two BE nodes

# Sample: imbalanced load

- ▶ Start a FE as seed on ecelinux1 taking 6 cores
- ▶ Start a BE on ecelinux2 taking 1 cores
- ▶ Start another BE on ecelinux2 taking 3 cores
- ▶ Start another BE on ecelinux5 taking 2 cores(Intel i7)
- ▶ Client send requests
- ▶ Request performance counters on the servers
- ▶ **Expecting:** BE with 3 cores receives roughly 3 times as much load as the BE with 1 core; BE on ecelinux5 with two cores receives roughly twice as much load as the BE with 1 core on ecelinux2, due to better hardware.
- ▶ Note: To reduce load on the ecelinux servers please keep your experiments short. Do not run the system at full throttle for more than about 10s at a time. This will help avoid interference between performance experiments executed concurrently by different groups.

# Sample: fault tolerance

- ▶ Start a FE as seed on ecelinux1 taking 4 cores
- ▶ Start a BE on ecelinux2 taking 2 cores
- ▶ Start another BE on ecelinux5 taking 2 cores
- ▶ Client send fixed number of requests to FE allowing both ecelinux2&5 have enough workload
- ▶ kill one of the BEs 5 seconds after start-up(requests have not finished then.)
- ▶ Request performance counters on the servers
- ▶ **Expecting:** client doesn't receive any exceptions and eventually the numRequestsCompleted at the FE equals the number of requests issued by the client.

# Use Piazza

► [piazza.com/uwaterloo.ca/summer2015/ece454750](https://piazza.com/uwaterloo.ca/summer2015/ece454750)

The screenshot shows the Piazza website interface for the ECE 454 / 750 course. The browser address bar displays the URL <https://piazza.com/class/i9ydvvke3m2c3?cid=36#>. The Piazza header includes the course name, navigation tabs (Q & A, Resources, Statistics, Manage Class), and the user's name (hua fan). Below the header, there are tabs for assignments (assignment1, assignment2, assignment3, midterm\_exam, final\_exam) and a 'Piazza Careers' link. The main content area is titled 'Class at a Glance' and shows a summary of the course's status: 36 total posts, 136 total contributions, 43 instructors' responses, 9 students' responses, and a 27 min avg. response time. A green bar indicates that 95 out of 133 students are enrolled. On the left side, there is a list of recent posts, including 'Connection Reset when accessing client', 'Server Load Info', 'Server arguments.', and 'Referencing jar files'.

**Class at a Glance** Updated 1 minute ago. [Reload](#)

- ✓ no unread posts
- ✓ no unanswered questions
- ✓ no unresolved followups

36 total posts  
136 total contributions  
43 instructors' responses  
9 students' responses  
27 min avg. response time

**Student Enrollment** ..out of 133 (estimated)

95 enrolled

**Filtering by: assignment1**

**TODAY**

- Connection Reset when accessing client** 10:46PM  
I am trying to connect a BEServer to a FEServer through the management interface... but every time I get org.apache.thr
- Server Load Info** 9:40PM  
I needed advice regarding measuring server load info. Lets say I have 1 FE and 3 BE's and a Parallel Client (All on
- Server arguments.** 5:52PM  
The server uses AbstractServerArgs as the argument . What are the subclasses of that thing ? I can just copy the tutoria
- Referencing jar files** 4:58PM  
Hi, If I reference only libthrift.jar then it is giving a build error. It seems I also need to reference all